

The learning process

Components of a trainable model

- x - input data (eg. an image containing blood cells, flower petal width/length)
- y - target data (eg. cell type classification, centroid of the flower clusters)
- θ - the learned parameters that characterize the model, f
- L - an objective function

$$\operatorname{argmin}_{\theta} L(y, f(x, \theta))$$

